



INSTREK TECHNOLOGIES

 **DOWNLOAD / SAVE PDF**

CONFIDENTIAL RISK ASSESSMENT REPORT

Your organisation carries *three open wounds* that your AI roadmap hasn't closed.

This report maps your organization's AI security exposure across the three pillars that regulators, auditors, and adversaries are already prioritising. It is based on your responses to the Instrek Technologies risk scan.

| PREPARED FOR               | ROLE                     | ORGANISATION | REPORT DATE   |
|----------------------------|--------------------------|--------------|---------------|
| Charusaraswat639649 Leader | CDO / Chief Data Officer | n            | 8 August 2026 |

56

RISK SCORE

Managed

Elevated

Critical

MODERATE RISK

Partial controls - governance gaps remain

## 01 - EXECUTIVE SUMMARY

# The situation your board will ask about next quarter.

Based on your self-assessment, your organization is operating with significant unmanaged AI risk. The gap is not primarily technical - it is architectural. Your teams have access to AI tools that your governance framework was not designed for, your agents operate without the identity and auditability of a regulated employee, and your cryptographic infrastructure is built for a threat landscape that has already shifted.

This is not a future risk. It is a present condition. The DPDP Act has given Indian regulators a mechanism to act. The NIST PQC standards have given your auditors a framework to measure against. And the velocity of agentic AI adoption in organizations means that every week without governance is a week of compounding exposure.

*"The organisations that define AI governance frameworks now will not just avoid the next breach. They will win the talent, the client trust, and the regulatory goodwill that defines who leads in the next decade of enterprise AI."*

- Instrek Technologies Research, April 2026

## 56

Your composite risk score out of 100. Above 70 indicates critical exposure requiring immediate architectural intervention.

### HIGHEST RISK PILLAR

Post-Quantum Defense - score 70/100. This is where your first 30 days of remediation should be focused.

### REGULATORY EXPOSURE

DPDP Act enforcement is active. Your current AI data flow posture may not withstand regulatory scrutiny without documented controls.

#### PEER CONTEXT

73% of organizations at your AI deployment stage have not completed a formal AI governance architecture. You are in the majority - and that is the risk.

## 02 - RISK PILLAR BREAKDOWN

# Where your exposure lives - and what it costs if unaddressed.

Your risk profile is mapped across the three AI security pillars that matter most to executives in 2026. Each pillar below reflects your specific answers, the gap they reveal, and the architectural response that closes it.

### AI Sovereignty & Data Leakage

Elevated Risk

*"If my India team uses ChatGPT to write code, where does my IP actually go? Can the vendor train on it?"*

Your assessment indicates partial or informal controls around GenAI tool usage in your India GCC. While some awareness exists, the absence of a formal audit trail means you cannot demonstrate to a regulator or auditor exactly what data has left your perimeter - or prove it hasn't. This is a documented but unquantified risk.

65

#### WHAT "CLOSING THIS GAP" LOOKS LIKE

A Private LLM deployment on your VPC or on-premise gives your developers full GenAI capability with zero data leaving your perimeter. Combined with a Secure RAG architecture, your proprietary codebases and internal documents become the model's knowledge base - not a training leak.

DPDP-compliant data flow mapping ensures every AI interaction is auditable and defensible to regulators.



## Agentic Accountability

Managed

*"What happens if an autonomous agent makes a \$1M financial error or creates a security backdoor - and I don't find out for weeks?"*

Your assessment indicates a mature approach to agentic AI governance. As your agent deployment scales, the priority is ensuring your framework grows with it - particularly as newer, more capable agents request broader system access. The edge cases that matter most are the ones your current framework was not designed for yet.

20

### WHAT "CLOSING THIS GAP" LOOKS LIKE

Treating AI agents as non-human employees - with their own credentials, scoped permissions, and immutable audit trails - converts an invisible risk into a manageable one. Human-in-the-loop (HITL) checkpoints at financial, security, and data-write actions ensure no agent can exceed a defined blast radius without human approval. An AI Governance framework gives your board a risk register entry with an owner.



## Post-Quantum & AI Cyber Defense

Critical Risk

*"Hackers are using AI to find vulnerabilities faster than we can patch them. And in five years, a quantum computer could decrypt everything we encrypted today."*

Your assessment indicates no formal cryptographic inventory has been conducted and post-quantum readiness is not on your active roadmap. Nation-state actors are currently executing "harvest now, decrypt later" operations - archiving your encrypted data today to decrypt when quantum capability arrives. NIST finalised its Post-Quantum Cryptographic standards in 2024. The migration timeline for large GCC environments is 3-5 years. The organisations that have not started are the ones that will not finish before the threat becomes real.

70

#### WHAT "CLOSING THIS GAP" LOOKS LIKE

A proactive, self-healing SOC uses AI to detect and auto-remediate threats before they reach human escalation - matching adversarial AI velocity with defensive AI velocity. A Post-Quantum Cryptography readiness assessment maps every RSA and ECC dependency in your environment and produces a NIST PQC-aligned migration roadmap. The organisations that start this migration now will complete it before the quantum threat is realised. Those that wait will not.

### 03 - REMEDIATION ROADMAP

## Three horizons. One coherent architecture.

Based on your risk profile, we recommend a structured remediation across three time horizons. This is not a linear checklist - these tracks run in parallel, governed by a single AI Security architecture that prevents the point-solution sprawl that most organizations fall into.

30d

IMMEDIATE - 0 TO 30 DAYS

### Contain the bleed: AI data flow audit + shadow AI policy

Before any architecture change, you need to know what you are dealing with. A rapid AI data flow audit maps every external LLM API call being made by your teams - including shadow usage.

Simultaneously, a published AI acceptable-use policy with enforcement creates a defensible posture for regulators while architectural controls are being built.

AI Data Flow Audit

Shadow AI Detection

Acceptable Use Policy

DPDP Gap Assessment

90d

SHORT-TERM - 30 TO 90 DAYS

### Build the foundation: Private LLM + Agentic Identity

Deploy a Private LLM environment in your VPC or on-premise - with Secure RAG indexing your internal knowledge bases. In parallel, implement agentic identity management - assigning non-human credentials, scoped access controls, and HITL audit checkpoints to every AI agent operating in production. This closes the two highest-velocity risk vectors simultaneously.

Private LLM Deployment

Secure RAG Architecture

Agentic Identity Framework

HITL Audit Trails

12m

STRATEGIC - 90 DAYS TO 12 MONTHS

## Future-proof: SOC AI uplift + PQC migration roadmap

With the immediate risks contained, the strategic horizon is about asymmetric capability. An AI-powered self-healing SOC shifts your security posture from reactive to proactive. A Post-Quantum Cryptography readiness assessment produces the migration roadmap your board will need - beginning the transition from RSA/ECC to NIST PQC-approved algorithms across your infrastructure before the regulatory mandate arrives.

AI-Powered SOC

Self-Healing Infrastructure

PQC Readiness Assessment

Crypto-Agile Migration

NIST PQC Roadmap

### 04 - PEER INTELLIGENCE

## Where executives at your level actually stand.

These benchmarks are drawn from Instrek Technologies's assessment data across 340 organisations scanned in Q1 2026. They give you an objective read on how your posture compares - and where the largest competitive differentiation opportunity sits.

#### ORGANIZATIONS WITH FORMAL AI GOVERNANCE

**27%**

Only 1 in 4 organizations has a documented AI governance framework with an assigned owner. The other 73% rely on informal policies that would not survive an audit.

#### ORGANIZATIONS WITH AGENTIC IDENTITY MGMT

**11%**

Fewer than 1 in 10 organizations deploying AI agents has implemented dedicated agent credentials. The rest operate with agents inheriting human access - invisible to audit systems.

#### ORGANIZATIONS THAT HAVE STARTED PQC PLANNING

**18%**

Post-quantum migration is a multi-year programme. The 18% that started in 2025-26 will complete ahead of the regulatory mandate. The majority will not.

AVG TIME TO DETECT SHADOW AI USAGE

94 days

The median organization takes over 90 days to detect that developers have introduced an unauthorised AI tool. In that window, the IP damage is already done.

| CAPABILITY                 | PEER AVERAGE             | YOUR STATUS                                 | MATURITY |
|----------------------------|--------------------------|---------------------------------------------|----------|
| AI Acceptable Use Policy   | 41% have a formal policy | Likely not yet - below peer average         |          |
| Private / On-Prem LLM      | 19% deployed             | Partial or exploratory                      |          |
| Agentic Identity Framework | 11% operational          | Operational - ahead of 89% of peers         |          |
| PQC Readiness Assessment   | 18% completed            | Not started - below the 18% that have begun |          |
| AI-Powered SOC             | 33% deployed or piloting | Likely not yet - majority position          |          |

YOUR NEXT STEP

This report surfaces the risk.  
*A conversation closes it.*

Our security architects review a maximum of 8 new organisations per month. A 30-minute strategy call is not a sales call - it is a working session

where we map your specific architecture against your highest-priority pillar and give you a concrete, actionable starting point. No commitment required.



### Book a 30-min Strategy Call

Live session with a security architect. We review your specific pillar profile and recommend a starting architecture. This week's slots are filling.



### Request a Written Architecture Blueprint

Prefer async? Email us with your report score and we will respond within 24 hours with a written architecture recommendation tailored to your organization size and sector.



### Download the Full AI Risk Framework

Our 40-page technical framework covers all three pillars with architecture diagrams, vendor evaluation criteria, and implementation checklists for each.



### Share This Report With Your Team

Forward this report link to your CISO, CTO, or technology head. A shared risk view accelerates internal alignment and shortens the path to a governance decision.

**BOOK MY STRATEGY CALL →**

⚠️ We reviewed 340 risk profiles in Q1 2026. The organisations that acted on their highest-risk pillar within 30 days of assessment reduced their measured exposure by an average of 61%. The ones that filed this report and moved on did not. The window to act before your next board AI risk review is now.



# 01

## **You book the call**

A 30-minute slot with a security architect. Bring your biggest question. We bring the answers.

# 02

## **We map your architecture**

We review your risk profile, your current stack, and your timeline. You leave with a concrete starting point - not a pitch deck.

# 03

## **You decide what happens next**

If there is a fit, we scope a pilot engagement. If there is not, we will tell you honestly and point you to the right resource.

Instrek Technologies - AI Security

[strategy@instrek.com](mailto:strategy@instrek.com)   [instrek.com](https://instrek.com)   [Privacy Policy](#)

This report is based on self-reported data from the Instrek Technologies risk assessment. It is intended as a strategic orientation tool and does not constitute a formal security audit. Risk scores are indicative and derived from heuristic models calibrated against peer data. Instrek Technologies accepts no liability for decisions made solely on the basis of this report without further technical due diligence.