

INSTREK MAGNETO

AI Readiness Assessment Platform

DEVELOPER TECHNICAL DOCUMENTATION

Scoring Logic · Architecture · Integration Guide · Report Engine

Document Version	v1.0 — Initial Release
Classification	Confidential — Developer Use Only
Product	Magneto AI Readiness Platform
Client	Instrek
Year	2026

This document provides complete technical specifications for implementing the Magneto AI Readiness Assessment platform, including all scoring logic, calculation engines, report generation rules, and integration patterns.

1. Platform Overview

1.1 What is Magneto?

Magneto is Instrek's enterprise AI Readiness Assessment platform. It evaluates an organisation's readiness to adopt, deploy, and scale Artificial Intelligence across 7 critical dimensions derived from Instrek's proprietary AI Readiness Questionnaire. The platform delivers a scored maturity report, risk quantification, opportunity identification, and a 100-day transformation roadmap.

1.2 Assessment Architecture

The platform follows a 5-stage flow:

1. Company Intake — Captures organisation profile (name, industry, revenue, size, role, AI investment level).
2. Assessment — 22 weighted questions across 7 dimensions. Each question has 5 scored options (1–5).
3. Email Gate — Corporate email verification before results are revealed.
4. Report Generation — All scoring, benchmarking, risk calculation, and roadmap logic executes client-side in JavaScript.
5. Results Dashboard — Interactive multi-section report with radar chart, dimension breakdown, risk cards, opportunity table, and roadmap.

The current prototype is a single-file HTML/JS/CSS application with no backend dependency. All logic runs in the browser. For production, the developer should migrate the assessment engine and report generation to a server-side API.

1.3 Technology Stack (Current Prototype)

Layer	Technology	Notes
Frontend	HTML5 / CSS3 / Vanilla JS	Single-file SPA. No framework.
Charts	Chart.js v4.4.1 (CDN)	Radar chart, gauge, bar charts.
Fonts	Google Fonts (CDN)	Syne, JetBrains Mono, DM Sans.
Data Store	Browser memory (JS objects)	No persistence. Session only.
Logo	Base64 embedded PNG	Instrek_Logo_White.png embedded.

2. Question Bank & Dimensions

2.1 The 7 Assessment Dimensions

All 22 questions are sourced directly from the Instrek AI Readiness Questionnaire (AI_Readiness_Questionnaire.xlsx). They are mapped to 7 dimensions as follows:

#	Dimension	ID (code)	Q Count	What it measures
01	Company Strategy	strategy	3	AI/ML strategy alignment, executive sponsorship, ROI-based use case prioritisation.

02	Business Functions	business	3	AI adoption across BUs, KPI definition, IT-business collaboration.
03	Data Readiness	data	4	Data centralisation, quality, governance, cataloguing, historical availability.
04	Technology	technology	3	Scalable infrastructure, ML tooling, data engineering, API integration layers.
05	Security & Governance	security	3	Data privacy, regulatory compliance (DPDP/GDPR), RBAC, ethical AI policies.
06	People & Talent	people	3	AI/ML expertise, employee AI awareness, change management for adoption.
07	Operations & MLOps	operations	3	MLOps practices, model monitoring, drift detection, AI workflow integration.
TOTAL			22	

2.2 Answer Scoring Scale

Every question has exactly 5 answer options scored on a fixed integer scale:

Score	Label	Meaning
1	Critical Gap	Not implemented. Organisation is completely exposed or unaware.
2	Early Stage	Awareness exists but no formal implementation. Ad-hoc or informal only.
3	Developing	Partial implementation. Covers some areas but significant gaps remain.
4	Capable	Mostly implemented with minor gaps. Systematic approach in place.
5	Leader / Best-in-Class	Fully implemented, continuously improved, and competitive differentiator.

3. Scoring Engine — Calculation Logic

3.1 Per-Question Score

Each question answered by the user stores a raw integer score (1–5) in a JavaScript answers object indexed by dimension and question index:

```
// answers[dimensionIndex][questionIndex] = score (integer 1-5)

answers = {
  0: { 0: 4, 1: 3, 2: 5 }, // strategy: Q1=4, Q2=3, Q3=5
  1: { 0: 2, 1: 3, 2: 2 }, // business
  2: { 0: 1, 1: 2, 2: 3, 3: 2 }, // data (4 questions)
  3: { 0: 3, 1: 4, 2: 3 }, // technology
  4: { 0: 2, 1: 3, 2: 2 }, // security
  5: { 0: 3, 1: 4, 2: 3 }, // people
  6: { 0: 2, 1: 3, 2: 3 }, // operations
}
```

3.2 Dimension Score

Each dimension is scored as a percentage of maximum possible points for that dimension:

```
function calcCatScore(catIndex) {
  const d = answers[catIndex];
  const pts = Object.values(d).reduce((acc, score) => acc + score, 0);
  const max = Object.keys(d).length * 5; // 5 = max score per question
  return max ? Math.round(pts / max * 100) : 0;
}

// Example — Data Readiness (4 questions, scores: 1,2,3,2)
// pts = 1+2+3+2 = 8
// max = 4 × 5 = 20
// score = round(8/20 × 100) = 40%
```

Note: Each dimension is weighted equally in the overall score. There is no per-question weighting in the current engine — the "weight" labels (High/Medium) on questions are displayed for respondent context only and do not affect the mathematical calculation. A weighted model is described in Section 3.6.

3.3 Overall Maturity Score (OMS)

The Overall Maturity Score is the simple arithmetic mean of all 7 dimension scores:

```
const dimScores = CATEGORIES.map((_, i) => calcCatScore(i));
// dimScores = [80, 47, 40, 60, 47, 67, 53] ← example

const overall = Math.round(
  dimScores.reduce((sum, s) => sum + s, 0) / CATEGORIES.length
);
// overall = round((80+47+40+60+47+67+53) / 7)
// overall = round(394 / 7)
// overall = round(56.28)
// overall = 56%
```

3.4 Maturity Tier Classification

The OMS maps to one of 5 named maturity tiers. Thresholds are fixed and not configurable via UI:

Tier Name	OMS Range	Display Colour	Description
Laggard	0 – 24%	#EF4444 (Red)	Critical foundational gaps. AI cannot deliver value at scale without urgent remediation.
Explorer	25 – 44%	#F59E0B (Amber)	AI journey started but critical gaps across multiple dimensions limit progress.
Builder	45 – 61%	#60A5FA (Blue)	
Scaler	62 – 77%	#2DD4BF (Teal)	Strong foundations and growing deployment. Focus on systematising and governing AI.
Leader	78 – 100%	#10B981 (Green)	Top-tier AI readiness. Challenge is compounding the advantage and deepening AI into the core business model.

```
function getTier(pct) {
  if (pct < 25) return { name: "Laggard", color: "var(--danger)" };
  if (pct < 45) return { name: "Explorer", color: "var(--warn)" };
  if (pct < 62) return { name: "Builder", color: "#60a5fa" };
  if (pct < 78) return { name: "Scaler", color: "var(--teal)" };
  return { name: "Leader", color: "var(--safe)" };
}
```

3.5 Industry Benchmark Comparison

Each dimension score is compared against industry benchmark averages. Benchmarks are hard-coded in the `INDUSTRY_BENCHMARKS` object and represent anonymised aggregate data from 340+ enterprise assessments (2025–2026).

```
const INDUSTRY_BENCHMARKS = {
  "BFSI": { strategy:55, business:45, data:40, technology:48, security:52,
  people:38, operations:42 },
  "Fintech": { strategy:62, business:52, data:55, technology:60, security:48,
  people:50, operations:56 },
  "Healthcare": { strategy:42, business:38, data:35, technology:40, security:50,
  people:30, operations:36 },
  "Retail & FMCG": { strategy:48, business:46, data:42, technology:45, security:35,
  people:38, operations:44 },
  "Manufacturing": { strategy:40, business:38, data:34, technology:42, security:36,
  people:30, operations:38 },
  "Technology": { strategy:70, business:62, data:68, technology:72, security:58,
  people:65, operations:66 },
  "Logistics": { strategy:44, business:40, data:36, technology:44, security:38,
  people:32, operations:42 },
  "Education": { strategy:35, business:32, data:28, technology:34, security:30,
  people:30, operations:28 },
  "Other": { strategy:45, business:40, data:38, technology:44, security:38,
  people:35, operations:40 },
};

// Benchmark delta displayed per dimension:
```

```
const diff = dimScore - benchmarkScore;
// Positive diff = above average (shown in green with + prefix)
// Negative diff = below average (shown in red)
```

3.6 Weighted Scoring (Recommended for Production)

The current prototype uses equal weighting across all questions and dimensions. For production, Instrek recommends implementing the following weighted model:

The "High Weight" and "Medium Weight" labels visible on questions in the assessment UI are for respondent context only. They do NOT currently affect the mathematical score. Implement the model below to activate true weighted scoring.

```
// Recommended question-level weight multipliers
const WEIGHT = {
  "High Weight": 1.5, // multiplied into score before aggregation
  "Medium Weight": 1.0,
  "Low Weight": 0.7,
};

// Weighted dimension score:
function calcWeightedScore(catIndex) {
  let weightedSum = 0;
  let weightedMax = 0;
  QUESTIONS[CATEGORIES[catIndex].id].forEach((q, qi) => {
    const w = WEIGHT[q.weight] || 1.0;
    const score = answers[catIndex][qi] || 0;
    weightedSum += score * w;
    weightedMax += 5 * w;
  });
  return weightedMax ? Math.round(weightedSum / weightedMax * 100) : 0;
}
```

4. Report Engine — Output Calculations

4.1 Risk & Financial Exposure

The report generates 6 risk/opportunity cards with estimated financial values. These are derived from the company revenue input and the Overall Maturity Score. The developer must note these are indicative estimates, not audited figures.

Risk Card	Formula	Logic
Regulatory / Compliance	$revMult \times (1 - OMS/100) \times 0.35$	35% of maximum exposure scales with maturity gap. 18% AI revenue uplift benchmark, scaled by gap.
Lost Revenue Upside	$revMult \times 0.18 \times (1 - OMS/100)$	

Operational Inefficiency	$revMult \times 0.12 \times (1 - OMS/100)$	12% productivity gain benchmark, scaled by gap.
Competitive Risk	"18–24 month lag" (qualitative)	Fixed qualitative label. No formula.
Talent Risk	"High risk" (qualitative)	Fixed qualitative label. No formula.
Agentic AI Gap	"Critical" (qualitative)	Fixed qualitative label. No formula.

Revenue multipliers (revMult) are mapped from the revenue band selected in the intake form:

```
const revMult = {
  "Under ₹50 Cr":      30,
  "₹50-250 Cr":        100,
  "₹250-1000 Cr":      300,
  "₹1000-5000 Cr":     1000,
  "₹5000 Cr+":         2500,
}[companyInfo.revenue] || 200;

// Example - Company with ₹500 Cr revenue, OMS = 40%:
// revMult = 300
// Penalty exposure = 300 × (1 - 0.40) × 0.35 = 300 × 0.60 × 0.35 = ₹63 Cr
// Lost revenue     = 300 × 0.18 × 0.60 = ₹32.4 Cr/yr
// Inefficiency     = 300 × 0.12 × 0.60 = ₹21.6 Cr/yr
```

4.2 Dimension Finding Text

Each dimension card in the report shows a qualitative finding paragraph. The text is selected from a pre-written findings library (CAT_FINDINGS object in JS) based on the dimension score tier:

```
function getCatFinding(catId, pct) {
  const f = CAT_FINDINGS[catId]; // object with .low, .mid, .high
  if (pct < 45) return f.low;     // Critical or Explorer tier
  if (pct < 70) return f.mid;     // Builder or lower Scaler
  return f.high;                 // Scaler or Leader
}

// Thresholds:
// pct < 45 → "low"  finding (critical gap narrative)
// 45-69   → "mid"  finding (developing narrative)
// 70+     → "high" finding (strength narrative)
```

4.3 Benchmark Delta Display

```
// For each dimension in the benchmark section:
const yours = catScores[i];           // user's dimension score (0-100)
const avg   = bench[cat.id];          // industry benchmark value
const diff  = yours - avg;             // delta

// Display rules:
// diff >= 0 → colour: green (#10B981), prefix: "+"
// diff < 0  → colour: red (#F43F5E),  prefix: "" (negative shown)
```

```
// Visual bar: two overlapping bars on same baseline
// "yours" bar fills to the user score %
// A vertical marker line sits at the benchmark %
```

4.4 Radar Chart Data

The radar chart plots two datasets simultaneously — the user's dimension scores and the industry benchmark scores:

```
// Chart.js radar config:
datasets: [
  {
    label: "Your Score",
    data: catScores,           // [80, 47, 40, 60, 47, 67, 53]
    borderColor: "#2563eb",
    backgroundColor: "rgba(37,99,235,0.12)",
  },
  {
    label: "Industry Avg",
    data: benchVals,           // from INDUSTRY_BENCHMARKS[industry]
    borderColor: "rgba(255,255,255,0.2)",
    borderDash: [6, 4],        // dashed line for benchmark
  }
]

// Scale: 0-100. Axis labels use dimension shortNames.
// Point colours: each point uses its dimension colour.
```

4.5 Score Gauge

The semicircular gauge is drawn on an HTML5 Canvas element. It uses three colour zones to convey maturity:

```
function drawGauge(pct, tierColor) {
  const cx = W/2, cy = H*0.75;
  const radius = Math.min(W, H) * 0.68;

  // Background arc (full semicircle, transparent grey)
  ctx.arc(cx, cy, radius, Math.PI, 2*Math.PI);

  // Three colour zone arcs (0-33% red, 33-62% amber, 62-100% blue):
  // [0, 0.33] → #EF4444 (danger)
  // [0.33, 0.62] → #F59E0B (warning)
  // [0.62, 1] → #2563EB (safe/blue)

  // Value arc: filled from 0 to (pct/100 × π radians)
  const endAngle = Math.PI + (pct/100) * Math.PI;
  ctx.strokeStyle = tierColor; // uses the tier's display colour
  ctx.shadowBlur = 12;
  ctx.shadowColor = tierColor; // glowing effect

  // Needle dot at arc endpoint
  ctx.arc(endX, endY, 8, 0, 2*Math.PI);
  ctx.fillStyle = tierColor;
}
```


4.6 100-Day Roadmap Logic

The roadmap is currently static — 3 fixed phases with pre-written action items. In the production build, these should be made dynamic based on the user's weakest dimensions:

Phase	Timeline	Logic for Dynamic Generation
Phase 1 — Foundation	Days 1–30	Inject action items targeting the 2–3 dimensions with scores below 40%. Prioritise Security and Strategy if they are in the bottom tier.
Phase 2 — Acceleration	Days 31–70	Focus on deploying AI use cases and closing gaps in Data and Technology dimensions. Inject dimension-specific actions for any dimension scoring 40–60%.
Phase 3 — Scale & Lead	Days 71–100	Board-level governance, scaling of successful pilots, and external positioning. Always appears in Phase 3 regardless of scores.

```
// Recommended dynamic roadmap trigger logic:
const criticalDims = CATEGORIES.filter((c, i) => catScores[i] < 40);
const developingDims = CATEGORIES.filter((c, i) => catScores[i] >= 40 && catScores[i] < 60);

// Phase 1 actions = criticalDims.map(d => ACTION_LIBRARY[d.id].urgent)
// Phase 2 actions = developingDims.map(d => ACTION_LIBRARY[d.id].accelerate)
// Phase 3 actions = always use scale/governance actions
```

5. Data Flow & State Management

5.1 State Object Schema

The assessment engine uses the following JS state objects. These should be migrated to a backend session or database in production:

```
// Company profile (captured on intake form)
companyInfo = {
  company: String,    // Free text company name
  industry: String,   // Selected from fixed list of 9 options
  revenue: String,    // Selected revenue band (5 options)
  size: String,       // Employee count band (5 options)
  role: String,       // Respondent role (7 options)
  invest: String,     // Current AI investment level (5 options)
}

// Assessment state
currentCat = Number;  // 0-6 (dimension index)
currentQ = Number;    // 0-N (question index within dimension)
```

```
// Answer store - populated as user answers questions
answers = {
  [catIndex: Number]: {
    [qIndex: Number]: Number // score: 1-5
  }
}

// Example populated answers object:
// { 0:{0:4,1:3,2:5}, 1:{0:2,1:3,2:2}, 2:{0:1,1:2,2:3,3:2}, ... }
```

5.2 Assessment Completion Check

```
// Total questions = CATEGORIES.reduce((s,c) => s + QUESTIONS[c.id].length, 0)
// = 3+3+4+3+3+3+3 = 22

// Assessment is complete when:
const totalAnswered = Object.values(answers)
  .reduce((sum, d) => sum + Object.keys(d).length, 0);

const isComplete = totalAnswered === 22;

// Navigation: after the final question (operations Q3), trigger showGate()
// Gate blocks results until corporate email is validated.
```

5.3 Email Gate Validation Logic

```
// Validation rules for email gate:
function isValidCorporateEmail(email) {
  const FREE_PROVIDERS = [
    "gmail", "yahoo", "hotmail", "outlook", "rediffmail", "yopmail"
  ];
  const isValidFormat = /^[^\s@]+@[^\s@]+\.[^\s@]+$/.test(email);
  const domain = email.split("@")[1]?.split(".")[0]?.toLowerCase();
  const isFreeProvider = FREE_PROVIDERS.includes(domain);
  return isValidFormat && !isFreeProvider;
}

// Rejection message for free providers:
// "Please enter a valid corporate email address (not Gmail/Yahoo)."
```

6. Production Architecture Recommendations

6.1 Recommended Tech Stack

Layer	Recommended	Purpose
Frontend	Next.js 14 / React	Component architecture, SSR for report pages, SEO for landing page.

Backend API	Node.js / FastAPI	Assessment session management, scoring engine, report generation.
Database	PostgreSQL	Store assessments, responses, leads, and benchmark data.
Auth / Gate	Custom OTP or Auth0	Email verification before report access. Lead capture.
PDF Export	Puppeteer / WeasyPrint	Generate downloadable branded PDF report.
Email	Resend / SendGrid	Send report link, follow-up sequence to leads.
CRM Integration	HubSpot / Salesforce API	Push lead data (email, company, score) to CRM on completion.
Analytics	Mixpanel / PostHog	Track assessment completion rates, drop-off per question, conversion.

6.2 Critical Production API Endpoints

```
// POST /api/assessment/start
// Body: { companyInfo }
// Returns: { sessionId: string }

// POST /api/assessment/answer
// Body: { sessionId, catIndex, qIndex, score }
// Returns: { progress, liveScore }

// POST /api/assessment/gate
// Body: { sessionId, email, name }
// Returns: { reportToken: string } OR { error: "invalid_email" }

// GET /api/report/:reportToken
// Returns: { answers, scores, report } - full computed report data

// POST /api/lead
// Body: { email, name, companyInfo, overallScore, tierName }
// Action: push to CRM, trigger email sequence
```

6.3 Database Schema (Minimal)

```
TABLE assessments
  id          UUID PRIMARY KEY
  session_id  VARCHAR(64) UNIQUE
  company_name VARCHAR(255)
  industry    VARCHAR(100)
  revenue_band VARCHAR(50)
  size_band   VARCHAR(50)
  respondent_role VARCHAR(100)
  invest_level VARCHAR(100)
  created_at  TIMESTAMP
  completed_at TIMESTAMP NULL
  email       VARCHAR(255) NULL -- set at gate
```

```

report_token VARCHAR(64) NULL -- set at gate

TABLE answers
  id          UUID PRIMARY KEY
  assessment_id UUID REFERENCES assessments(id)
  cat_index   SMALLINT -- 0-6
  question_index SMALLINT -- 0-N
  score       SMALLINT -- 1-5
  answered_at TIMESTAMP

TABLE dimension_scores -- computed on completion
  assessment_id UUID REFERENCES assessments(id)
  dimension_id  VARCHAR(20) -- strategy, business, data...
  score_pct     SMALLINT -- 0-100
  overall_pct   SMALLINT -- 0-100
  tier          VARCHAR(20) -- Laggard|Explorer|Builder|Scaler|Leader

```

7. UI Component Reference

7.1 Live Seismograph

The animated seismograph is a Canvas element that visually represents AI readiness signal in real time as questions are answered. It is cosmetic but provides strong engagement value.

```

// Risk level fed to seismograph = 1 - (currentScore / maxPossibleScore)
// Risk 0.0 = perfect score → flat green line
// Risk 1.0 = zero score → highly erratic red line

// Colour mapping:
// risk < 0.3 → #10b981 (green) -- Strong signal
// risk < 0.6 → #f59e0b (amber) -- Developing signal
// risk >= 0.6 → #f43f5e (red) -- Critical gaps detected

// Amplitude: amp = risk × canvasHeight × 0.38
// Each frame: v = center - (random × amp + sin(time) × amp × 0.5)
// Data is a rolling buffer equal to canvas width in pixels.
// Each frame: shift left, push new value, redraw full line.

```

7.2 Category Navigation Strip

The category nav strip inside the assessment panel shows progress across all 7 dimensions. A dimension nav button shows three states:

- Default (grey) — dimension not yet reached.
- Active (blue underline) — current dimension being answered.
- Done (green + ✓ prefix) — all questions in this dimension answered.

Clicking a completed or current dimension jumps directly to its first unanswered question.

7.3 Progress Bar

The linear progress bar at the top of the assessment panel fills based on total questions answered out of 22:

```
const totalAnswered = Object.values(answers)
  .reduce((s,d) => s + Object.keys(d).length, 0);
progressFill.style.width = (totalAnswered / 22 * 100) + "%";
```

7.4 Live Score Counter

The four live metric cells at the top of the assessment panel update on every answer:

- "Total Score" — raw sum of all answer scores so far.
- "Maturity %" — $(\text{totalScore} / (\text{answeredCount} \times 5)) \times 100$, colour-coded by tier.
- "Dimension" — "currentCat+1 / 7" counter.
- "Remaining" — 22 minus totalAnswered.

8. Extending the Platform

8.1 Adding a New Industry to Benchmarks

- Add a new key to INDUSTRY_BENCHMARKS with 7 dimension scores (0–100 each).
- Add the industry name to the intake form <select> options.
- The report engine will automatically use the new benchmark values.

8.2 Adding Questions to a Dimension

- Add the question object to QUESTIONS[dimensionId] array.
- Update the total question count in all references (currently hardcoded as 22 in some display strings).
- No changes to the scoring engine are needed — it dynamically counts questions per dimension.

If you add questions, update the stats strip on the landing page ("22" → new count) and the gate description text ("22 questions").

8.3 Adding a New Dimension

1. Add a new entry to the CATEGORIES array with id, name, icon, color, shortName, desc.
2. Add the question bank: QUESTIONS[newId] = [...questions...]
3. Add benchmark values for all industries: INDUSTRY_BENCHMARKS[each].newId = value
4. Add dimension findings: CAT_FINDINGS[newId] = { low, mid, high }

- 5. Update the OMS divisor (currently 7 = CATEGORIES.length — handled automatically).

8.4 PDF Report Export

The current version has no PDF export. For production, implement a server-side PDF generator:

```
// Recommended approach: Puppeteer (headless Chrome)
// 1. User clicks "Download PDF" on results page
// 2. POST /api/report/pdf with { reportToken }
// 3. Server loads /report/:token in headless Chrome
// 4. Puppeteer renders page, captures as PDF
// 5. Apply Instrek branded header/footer via PDF options
// 6. Return PDF blob for download

// Alternative: Generate HTML report string server-side,
// render with WeasyPrint for Python-native PDF output.
```

8.5 CRM Lead Push

```
// Trigger on gate submission (after email validation):
await fetch("/api/lead", {
  method: "POST",
  body: JSON.stringify({
    email: gateEmail,
    name: gateName,
    company: companyInfo.company,
    industry: companyInfo.industry,
    revenue: companyInfo.revenue,
    role: companyInfo.role,
    overallScore: overall,
    tier: tier.name,
    weakestDim: weakest.name,
    strongestDim: strongest.name,
    dimScores: catScores,
  })
});

// CRM field mapping (HubSpot example):
// email → Contact Email
// overallScore → Custom: AI Maturity Score
// tier → Custom: AI Maturity Tier
// industry → Industry
// company → Company Name
```

9. Brand & Design Tokens

9.1 Colour Palette

Token	Hex Value	Usage
--bg	#03071E	Page background — deep navy.

--bg2	#060D2E	Cards and panel backgrounds.
--bg3	#0A1535	Elevated panels (panel header, seismo bar).
--accent (primary)	#2563EB	All primary CTAs, active states, highlights.
--accent2	#3B82F6	Hover states and gradient endpoints.
--text	#EEEEEE6	Primary body text — warm off-white.
--muted	#5A6272	Secondary text, labels, placeholders.
--danger	#F43F5E	Error states, low scores, risk indicators.
--warn	#F59E0B	Warning states, medium scores.
--safe	#10B981	Success states, high scores.
Instrek Orange	#E8602C	Instrek brand accent — used on logo only in Magneto. Not used as UI accent.

9.2 Typography

Font Family	Weights	Usage
Syne	400, 600, 700, 800	All headings, hero text, category names, score numbers, buttons.
JetBrains Mono	300, 400, 500	Score numbers, labels, tags, eyebrows, code references.
DM Sans	300, 400, 500	Body text, form inputs, answer option text.

9.3 Logo Usage

The Instrek logo (Instrek_Logo_White.png) is embedded as a base64 data URI in the HTML file. It appears in:

- Navigation bar (height: 32px)
- Email gate modal (height: 28px)
- Results report navigation (height: 28px)
- Sprint CTA card (height: 20px)
- Footer (default height)

For production: serve the logo from a CDN/static asset server instead of embedding it as base64. The base64 encoding adds ~1.8MB to the HTML file size.

10. Changelog & Roadmap

10.1 Version History

Version	Date	Changes
v0.1	Apr 2026	Initial prototype — 3 industry verticals, 12 questions, basic scoring.
v0.2	Apr 2026	Expanded to 5 industries, 48 questions across 4 pillars, answer options 1–4.
v0.3	Apr 2026	Added benchmark dashboard, ROI calculator, testimonials, deliverables sections to landing page.
v0.4	Apr 2026	Rebuilt as focused AI Readiness tool — 6 dimensions, 48 questions, full report engine.
v1.0	Apr 2026	Final prototype — Instrek questionnaire integrated (7 dimensions, 22 questions), Instrek logo embedded, LinkedIn blue palette applied. This documentation version.

10.2 Recommended Production Roadmap

Priority	Feature	Notes
P0 — Critical	Backend API + Database	Replace browser-only state. Enable lead capture, persistence, analytics.
P0 — Critical	Email gate with OTP verification	Replace basic validation with real email verification.
P1 — High	PDF Report Export	Puppeteer-based branded PDF for download and sharing.
P1 — High	CRM Integration	Auto-push leads to HubSpot/Salesforce on report unlock.
P2 — Medium	Weighted Scoring Engine	Activate question-level weight multipliers (Section 3.6).
P2 — Medium	Dynamic Roadmap Generation	Generate roadmap actions from weakest dimensions (Section 4.6).
P2 — Medium	Admin Dashboard	Instrek internal view of all assessments, leads, scores, and trends.
P3 — Nice	LinkedIn Share Card	OG image auto-generated with company name + score for social sharing.
P3 — Nice	Benchmark Data Refresh Pipeline	Auto-update INDUSTRY_BENCHMARKS from real assessment data monthly.

This document is maintained by Instrek. For questions on implementation, contact the Instrek product team. All scoring logic described here reflects the v1.0 prototype. Production implementations may deviate from this specification — ensure this document is updated to reflect any such changes.